AI BUSINESS EXPERT

Intelligent Automation Is Delivering

But Only Where the Work Was Done First

20 March 2026

Internal Research Brief

Internal Research Brief — AI Business Expert Date: 20 March 2026 | Expanded Edition

---

Why This Matters Right Now

The debate about whether intelligent automation can deliver at scale is over. Not because the technology became simpler, or the risks disappeared, but because enough named enterprises with disclosed financials have now run these programmes long enough to report results that hold up to scrutiny. The question has moved — from “will this work?” to “why isn’t it working here, and what did the organisations where it is working do differently?”

JPMorgan has been processing 12,000 commercial loan agreements annually through its COiN (Contract Intelligence) system since 2017 — work that previously consumed 360,000 legal and loan officer hours per year (Bloomberg, 2017). That is nearly a decade of verified operational reality, not a pilot result or a vendor projection. Aviva documented £60 million in insurer savings from deploying 80 AI models across its claims function, with the figures disclosed to investors — the highest-confidence single-organisation ROI figure in the public evidence base (McKinsey & Company, 2025). IBM ran itself as the test environment, deploying 155 AI use cases across its own 175-country operation before selling the approach to clients, and now states US$4.5 billion in annual run-rate productivity savings, with 3.9 million hours saved in 2024 alone (IBM, 2025; Benzinga, 2025).

These are not edge cases. They represent a coherent pattern: enterprises that committed to process redesign before technology deployment, that chose high-volume repetitive workflows as first targets, and that invested in measurement from day one. The organisations that did not do those things — and the research

is explicit that the majority did not — are sitting on automation programmes that deliver narrow results, generate high maintenance overhead, and rarely expand beyond their pilot boundaries.

The urgency is structural. Deloitte's 2026 Technology Trends study found that 60% of workers are now equipped with AI tools, up from 40% just one year earlier — a pace of change that is compressing the window between early movers and the field (Deloitte, 2025). Gartner projects that 40% of enterprise applications will feature embedded AI agents by end of 2026, up from less than 5% in 2025 (Gartner, 2025). This is not a market forecast. It is a product roadmap signal from Salesforce, ServiceNow, SAP, and Microsoft — platforms organisations already own — that capability will expand materially whether or not internal teams are ready to use it.

The cost of waiting is not static. It is a compounding disadvantage.

The State of Play: What the Data Actually Shows

Three independent research programmes published in late 2025 converge on the same structural conclusion: AI adoption is broad but value is concentrated in a small cohort of organisations that approach deployment differently.

McKinsey's November 2025 State of AI survey, covering 1,993 respondents across industries and geographies, found that 88% of organisations deploy AI in at least one business function — a near-universal adoption rate. Yet only 6% can attribute more than 5% of EBIT to AI (McKinsey & Company, 2025, November). The gap between deployment and value is not a technology problem. It is a process and measurement problem.

Deloitte's 2026 research adds a human dimension: 25% of leaders now report transformative AI effects — a figure that doubled year-over-year — but 36% of organisations expect more than 10% of jobs to be fully automated within a single year, creating a change management challenge that most automation roadmaps treat as secondary (Deloitte, 2025). Organisations that treat workforce transition as an afterthought are discovering that the change management failure is what stops programmes from scaling, not technology inadequacy.

Forrester's Predictions 2026 report identifies a specific inflection risk: fewer than 15% of firms will activate the agentic automation features now embedded in their platforms by end of 2026 — not because they lack access, but because they lack the internal process clarity and governance structures required to deploy them safely (Forrester Research, 2025). Platform investment is not automatically converted into operational capability.

Ernst & Young's research on robotic process automation programmes found that up to 50% of initial RPA projects fail — and the causes are consistent: wrong process selection, insufficient change management, and underestimating the fragility of UI-dependent bots when application interfaces change (EY, 2023). The lesson is not that RPA does not work. The lesson is that automation deployed on broken or poorly understood processes automates the dysfunction.

The evidence base is now sufficiently large to draw clear conclusions: the organisations achieving material returns are not using different technology. They are using the same technology differently — with more rigour on process selection, more investment in measurement, and a stronger organisational commitment to redesign before deployment.

The Problem Before: Why Automation Has Historically Been Hard

Understanding why intelligent automation is succeeding where it is — and failing where it is — requires understanding what made it difficult in the first place.

For most enterprises, the first wave of automation arrived in the form of robotic process automation: software bots that mimicked human interactions with user interfaces, navigating screens, extracting data, entering values, and triggering downstream processes. The economics were appealing. The reality was fragile. Every time an application updated its interface, bots broke. Every time a process owner changed a workflow step, maintenance was required. IT teams found themselves managing increasingly complex bot estates with mounting technical debt, while process owners waited weeks for fixes that once took hours to complete manually.

The second wave introduced cognitive capabilities — OCR, natural language processing, machine learning-based classification — but brought a different problem: model governance. Who owns the model? Who retrains it when accuracy drifts? Who is accountable when the system makes a decision that causes a customer complaint or a compliance breach? Most organisations lacked the data science capacity to answer those questions in production, and many pilots failed to survive contact with operational reality.

The third problem was process quality. Automation vendors discovered — often after contract signature — that the processes their clients wanted to automate were not actually well-defined. They were handled differently by different staff, contained undocumented exception paths, and relied on tacit knowledge that no

one had bothered to write down because experienced humans could handle the variation intuitively. You cannot automate what you cannot specify.

The fourth problem was measurement. Automation programmes were initiated with business cases built on projected FTE savings, but follow-through on measurement was inconsistent. Programmes that could not demonstrate ROI in concrete terms were vulnerable to budget cuts when economic conditions tightened — and many were cut regardless of whether they were actually working.

Intelligent automation in 2025 and 2026 is working in the organisations that solved these four problems first: they chose processes that were well-defined and high-volume; they invested in governance structures before deployment; they redesigned workflows rather than automating existing ones; and they built measurement frameworks into programme design from the start. None of these are technology capabilities. All of them are organisational disciplines.

Financial Services: Where the Evidence Is Strongest

Financial services has the deepest and most independently verifiable automation evidence base of any industry — partly because the regulatory environment demands that financial impacts be disclosed, and partly because the volume and repetitiveness of financial workflows make them structurally well-suited to automation.

JPMorgan — Contract Intelligence (COiN)

JPMorgan's COiN system processes 12,000 commercial loan agreements annually in seconds. The same work previously required 360,000 hours per year of lawyer and loan officer time — approximately 173 full-time equivalents working exclusively on document review for one product category (Bloomberg, 2017; klover.ai, 2025). COiN uses a combination of natural language processing and machine learning to extract key clauses, flag non-standard terms, and route documents for exception handling. The system has been in production for nine years — a proof of durability that many newer automation claims cannot match.

The legacy problem COiN solved was not just speed. Manual contract review was inconsistent: different reviewers applied different interpretations to ambiguous clauses, creating compliance exposure that was difficult to audit. The automation created a consistent, auditable review trail — a quality improvement that the hours-saved figure alone does not capture.

Commercial legal AI tools — including Harvey, Luminance, and Litera Kira — are now approaching comparable accuracy on standard contract types without requiring JPMorgan's proprietary fine-tuning investment, making the same

workflow transformation accessible to organisations without a global bank's AI budget.

Aviva — Claims Function Transformation

Aviva's automation programme, implemented in partnership with McKinsey QuantumBlack, deployed 80 AI models across its claims liability assessment function. The problems it addressed were specific and well-documented before the programme began: complex liability cases were taking significantly longer than necessary because routing decisions were made manually by adjusters with inconsistent expertise; claims were frequently routed to the wrong specialist, requiring re-routing and restarting assessment; and customer complaints about claims handling were high relative to industry peers.

After deployment: complex liability assessment time fell by 23 days per case. Routing accuracy improved by 30%. Customer complaints fell by 65%. The documented financial savings of £60 million (approximately AU$115 million) are investor-disclosed — the highest-confidence single-organisation ROI figure in the public evidence base for financial services automation (McKinsey & Company, 2025).

The programme's success is attributed specifically to the scope of process redesign that preceded model deployment. McKinsey QuantumBlack did not automate the existing claims process. It redesigned the triage and routing logic before building models on top of the new design.

Revenue Cycle Management in Healthcare Finance

Healthcare finance represents one of the clearest automation opportunity areas across any industry, and the results in 2025 reflect significant progress on a long-standing problem. Hospital revenue cycle management — the process of converting clinical activity into cash collections, covering patient registration,

insurance verification, coding, claims submission, denial management, and collections — has historically been characterised by error rates, processing delays, and administrative costs that consume 25–30% of total hospital revenue (Council for Affordable Quality Healthcare, 2024).

Methodist Health System achieved 71% of accounts removed from staff queues through AKASA's AI-powered RCM automation, resolving 56,118 accounts in eight months and saving 5,559 hours of staff time (AKASA, 2024). One healthcare provider achieved a 30% reduction in accounts receivable days within a single quarter, unlocking US$2.5 million in accelerated cash flow. Clean claim rates — the percentage of claims submitted without errors on first pass — improved from 92% to 99% at organisations deploying AI-assisted coding verification, with coding staff productivity doubling (Notable Health, 2025).

The industry-level prize is substantial. The Council for Affordable Quality Healthcare estimates that automating administrative RCM processes could save the US healthcare system US$9.8 billion annually — value that currently sits in processing inefficiency rather than patient care.

Manufacturing & Supply Chain: From Downtime to Data

Manufacturing and supply chain represent the second major evidence cluster — both in predictive maintenance, where sensor data and machine learning have a decade of proven application, and in supply chain optimisation, where AI is creating value at a scale that older planning systems could not reach.

Siemens — Predictive Maintenance and Procurement

Siemens has deployed AI-driven predictive maintenance across its manufacturing operations, using continuous sensor monitoring and machine learning to flag equipment failure risk before breakdown occurs. The business outcomes are among the most well-documented in manufacturing: a 50% reduction in unplanned line stoppages, a 45% reduction in unplanned downtime, and a 28% reduction in maintenance spend (Industrial Equipment News, 2024). These are not pilot results from a single facility — they represent outcomes from Siemens' broader manufacturing network.

The problem Siemens solved is one most manufacturers recognise: unplanned downtime is expensive not only because of the direct cost of the failure event, but because it disrupts production schedules, creates quality risks in work-in-progress at the time of failure, and generates urgent labour costs at premium rates. Preventive maintenance on fixed schedules replaced components that did not need replacing while missing failures that occurred between intervals. Predictive maintenance eliminates both categories of waste.

Siemens also deployed AI in procurement — an application that receives less attention than predictive maintenance but has produced comparable returns. The

AI procurement system evaluates suppliers across 200+ criteria, covering quality performance, delivery reliability, financial stability, ESG compliance, and geopolitical risk exposure. Processing 15,000+ active suppliers with this level of analytical depth would be impossible with manual review. The result: a 60% reduction in procurement cycle time and an 11% reduction in procurement costs (Siemens Digital Industries, 2025).

UPS — ORION Route Optimisation

UPS's ORION (On-Road Integrated Optimisation and Navigation) system is one of the oldest and most financially material logistics automation programmes in the public record. By applying machine learning to route optimisation for its parcel delivery network — minimising driving distance while accommodating delivery time windows, vehicle capacity, and real-time traffic conditions — UPS estimates annual savings exceeding US$400 million (UPS, 2023). Each mile eliminated from a driver's daily route across a fleet of hundreds of thousands of vehicles compounds to a saving that no manual planning process could approximate.

The operational problem ORION solved was the combinatorial complexity of route planning at scale. A driver with 100 stops on their route faces more possible sequences than can be evaluated by any human planner. Heuristic approaches left significant distance — and therefore fuel cost and driver time — on the table. The AI system does not find the perfect route. It finds a materially better one, consistently, across every driver and every day.

Nike — Supply Chain Intelligence

Nike's AI-powered supply chain programme covers more than 500 manufacturing and logistics facilities across 40+ countries, managing over 120,000 active SKUs. Before the programme, the company relied on traditional demand forecasting approaches that accumulated error at each node of the supply chain, resulting in inventory imbalances — excess stock in some markets, stockouts in others — that

damaged both margins and customer experience. AI-driven demand forecasting and inventory optimisation reduced lead times by 50% and achieved 99.7% fulfilment accuracy while cutting excess inventory costs by 23% (Nike, 2024). Working capital freed from inventory reduction flows directly to the balance sheet.

IT & Enterprise Operations: The Self-Service Revolution

IT service management and enterprise operations represent the application area with the most directly attributable labour cost savings in the public evidence base — because IT support volumes are large, measurable, and pre/post comparisons are straightforward.

IBM — Client Zero

IBM's Client Zero programme is the most comprehensive self-automation case study available. IBM deployed 155 AI use cases across HR, IT, procurement, finance, and sales within its own organisation before offering the approach commercially. The results across individual use cases are specific and consistently material.

IBM AskIT, an AI-powered IT support assistant, reduced standard IT support tickets by 56% between 2022 and 2024. By March 2025, 86% of IT queries were being resolved without human involvement — a self-service rate that would have been inconceivable at enterprise scale five years earlier (IBM, 2025). IBM AskHR resolves 94% of HR policy, benefits, and payroll queries autonomously, cutting HR support ticket volume by 75% and redirecting HR staff capacity toward advisory and strategic work. IBM states US$4.5 billion in annual run-rate productivity savings from these programmes — IBM-stated rather than independently audited, but consistent with 3.9 million hours saved in 2024 alone (Benzinga, 2025).

The problem IBM solved is the same problem every large organisation faces: a disproportionate volume of employee time is consumed by routine, repetitive enquiries that require access to information rather than judgement. Policy

questions, status checks, standard troubleshooting steps, password resets, leave balance requests, benefits explanations. These are answerable by well-designed AI systems. Humans are expensive and overqualified for this work.

BT Group — AIOps and Self-Healing Infrastructure

BT Group deployed Dynatrace AIOps across its Openreach infrastructure operations, targeting the specific problem of reactive IT maintenance: field engineers were being dispatched to address faults that either had not yet fully materialised, or that could have been resolved remotely with earlier detection. The AIOps deployment enables predictive fault identification and automated remediation without human dispatch in a growing proportion of cases. Results: a 10% reduction in repair volumes, a 4% reduction in energy consumption across managed infrastructure, and a projected cumulative saving of £28 million by 2027 (BT Group, 2024). The programme is on track and is being extended to additional network infrastructure categories.

Microsoft 365 Copilot — Enterprise Productivity at Scale

Forrester's Total Economic Impact study on Microsoft 365 Copilot modelled a composite enterprise of 3,250 users and found a 116% ROI over three years, with full payback in under six months (Forrester Research, 2024, October). Seventy percent of users surveyed reported higher productivity, with task completion time falling by 29% on average. Two named deployments illustrate the range of application.

Commercial Bank of Dubai deployed Microsoft 365 Copilot across its operations and saved 39,000 staff hours annually — hours previously spent on meeting summarisation, document drafting, data analysis, and report generation (Microsoft, 2025). British Columbia Investment Corporation (BCIC) saved more than 2,300 hours through Copilot-assisted investment research and reporting (Microsoft, 2025). Both represent deployments in knowledge-worker

environments where the value driver is cognitive task acceleration rather than process automation in the traditional sense.

Customer Experience: From Queue to Resolution

Customer-facing automation represents a different value proposition from back-office process automation. The primary metric is not hours saved — it is customer experience quality, expressed through handle time, abandonment rate, first-contact resolution, and complaint volume.

Iron Mountain — Salesforce Einstein Service Cloud

Iron Mountain, the information management and storage company, deployed Salesforce Einstein Service Cloud to address a specific and measurable operational problem: agents were navigating between multiple applications to gather customer context before they could respond to each enquiry. The average handle time was high relative to industry benchmarks, and chat abandonment rates were elevated because customers were waiting too long for substantive responses.

After deploying Einstein AI: 80% of cases were closed using AI-generated responses. Of those, 76% were sent by agents without any edits — a proxy metric for the quality and relevance of the AI-generated content. Chat abandonment fell from 5% to 1.5% (Salesforce, 2024). The metrics are from a Salesforce-published case study and should be read as indicative, but the operational mechanism they describe is consistent with comparable deployments in similar service environments.

Bank of America — Erica

Bank of America's virtual financial assistant Erica has now surpassed 3 billion client interactions since launch — a scale that makes it one of the most

extensively deployed conversational AI systems in financial services globally (Bank of America, 2025, August). Internally, Bank of America deployed Erica's underlying technology to its 213,000 employees as a productivity assistant, cutting IT service desk calls by more than 50% and delivering productivity gains of more than 20% for software developers specifically (Bank of America, 2025, April). The developer productivity finding is consistent with Deloitte research showing that AI copilots used alongside humans deliver 1.6x higher productivity growth than automation-alone approaches (Deloitte, 2025).

Healthcare: Quiet Wins in a Constrained System

Healthcare presents a unique automation context. The potential returns are large — administrative costs are estimated at 25–30% of total hospital spend — but the risk environment is more constrained than most industries, and the change management complexity in clinical settings is higher. The result is that healthcare automation wins in 2025 are concentrated in administrative and operational functions rather than clinical decision-making.

Clinical Documentation — AI Ambient Scribing

One of the highest-value near-term applications in healthcare is AI ambient scribing: the use of AI to listen to clinician-patient conversations and generate structured clinical notes automatically. The problem it solves is concrete and widely reported. Clinicians in the United States spend an estimated two hours on electronic health record documentation for every one hour of direct patient care — a ratio that drives burnout, reduces patient throughput, and contributes to early career exit.

Healthcare organisations using Commure Ambient AI reported an average 81% reduction in documentation time, with clinical notes closed in an average of 43 seconds after appointment completion and claim denial rates falling by more than 25% (Commure, 2025). The denial reduction finding is particularly important: poor clinical documentation is a primary driver of insurance claim denials, meaning documentation quality has a direct cash flow impact that extends well beyond clinician productivity.

Revenue Cycle — Administrative Automation

Beyond the Methodist Health and individual provider results cited earlier, the systemic shift is visible in survey data. Eighty per cent of health systems are now exploring, piloting, or implementing AI tools for revenue cycle management — a 38% increase in less than two years (KLAS Research, 2025). The pipeline of organisations entering production deployment suggests that the case study base will expand significantly through 2026.

Home Care Delivered achieved a 95% reduction in claims processing time for secondary insurance transfers through RPA, saving 416 staff hours annually on a single workflow — a result that illustrates both the potential of targeted automation and the realistic scale of single-workflow programmes (Accelirate, 2025). The honest implication: individual workflow automation wins are real but bounded. Portfolio-level value requires multiple programmes running simultaneously.

What Separates Winners from the 94%

McKinsey's most important finding from the November 2025 State of AI survey is not about technology. It is about what the 6% of organisations achieving material EBIT impact from AI did that others did not. The single strongest predictor of AI programme value is process redesign before deployment. Fifty-five per cent of AI high performers fundamentally redesigned underlying processes before deploying automation — nearly three times the rate of other organisations (McKinsey & Company, 2025, November).

Process redesign is not a technology activity. It requires operations leadership to map existing processes at a granular level, identify where value is being destroyed by inefficiency, exception handling, or unnecessary handoffs, and redesign the process to be automation-ready before the automation is built. This work takes time — typically three to six months for a meaningful process redesign effort — and it requires skills that are not in most IT or digital transformation teams.

Deloitte's 2026 research adds a complementary finding: 85% of organisations deploying AI agents are customising them with proprietary data, workflows, and integrations (Deloitte, 2025). Generic deployment of out-of-the-box AI tools delivers generic results. The organisations achieving material returns are treating AI deployment as an engineering activity that requires adaptation to their specific operational context, not a software procurement activity that delivers value upon installation.

Three other consistent characteristics of high-performing automation programmes appear across the research base:

Measurement discipline: Gartner found that fewer than 20% of companies excel at measuring hyperautomation effectiveness (Gartner, 2024). Programmes that

cannot demonstrate ROI in concrete terms are vulnerable to budget cuts and cannot make evidence-based decisions about where to scale or retire investment. Measurement frameworks need to be designed before deployment — not retrofitted after questions arise.

Right process selection: The EY finding that up to 50% of initial RPA projects fail is concentrated in organisations that chose processes because they seemed automatable, not because they had the characteristics that predict automation success: high volume, low exception rate, stable and documented logic, and a data trail that enables performance measurement. Process selection criteria are as important as technology selection criteria.

Change management investment: Automation programmes that treat workforce transition as an afterthought discover that the human resistance — not the technology — is what limits scale. Employees whose roles change significantly due to automation require retraining, role redefinition, and active communication about what the programme means for their career path. Organisations that invest in this consistently report faster programme acceptance and higher quality outcomes, because staff actively engaged with the automation contribute to improving it.

The Measurement Problem

There is a measurement failure sitting inside most automation programmes that is rarely discussed openly. Organisations launch programmes, report headline savings, and move on — without building the feedback loops that would allow them to know whether the automation is still working twelve months later, whether its performance has drifted, or whether a changing process environment has made it obsolete.

Gartner's finding that fewer than 20% of companies excel at measuring hyperautomation effectiveness is not a finding about data capability. Most organisations have the data they would need to measure automation performance. It is a finding about governance: no one has been made responsible for ongoing measurement, and without that accountability, the reporting does not happen (Gartner, 2024).

The practical consequence is visible in bot estates. In organisations with mature RPA programmes, it is common to find a significant proportion of bots running processes that have changed around them — processing data that is no longer accurate, flagging exceptions that are no longer exceptions, or silently producing outputs that no one is checking. The cost is not just the direct error cost. It is the opportunity cost of investment that could have been redirected to higher-value automation.

The measurement standard that McKinsey identifies in high-performing programmes includes three elements: a defined baseline captured before deployment (so improvement can be measured against it, not against a projection); an automated monitoring mechanism that flags performance deviation without requiring manual investigation; and a governance cadence —

typically quarterly — where programme owners review measurement outputs and make explicit decisions about continuation, scaling, or retirement.

The Agentic Horizon: What Changes Next

The next phase of intelligent automation is agentic: AI systems that do not simply process a defined input and produce a defined output, but that take multi-step action sequences, use tools, make intermediate decisions, and complete tasks that previously required sustained human attention over minutes or hours.

The capabilities are real. The commercial deployments at scale are still limited. Gartner projects that more than 40% of agentic AI projects will be cancelled by the end of 2027 — not because the technology does not work, but because organisations are deploying agents into processes that are not well enough defined to benefit from autonomous action (Gartner, 2025). The same process quality problem that limited first-generation RPA is creating the same ceiling for agentic AI.

Deloitte describes the current practical guidance as the “agentish” path: AI-assisted deterministic workflows where AI contributes intelligence — summarisation, classification, drafting, recommendation — within a human-supervised process, rather than full autonomy. This approach captures a meaningful portion of agentic AI’s value while carrying materially lower risk (Deloitte, 2025). For most organisations, this is the right near-term frame.

The Gartner projection for 2026 — 40% of enterprise applications featuring embedded AI agents — should be understood not as a deployment target but as a capability signal. Salesforce Agentforce, ServiceNow Now Assist, Microsoft Copilot Studio, and SAP Joule are shipping agentic capabilities into platforms organisations already own. The automation landscape of 2027 will look materially

different from today not because organisations will have built new systems, but because the systems they already pay for will have changed around them.

The Risk Side of the Ledger

A credible account of intelligent automation's current state requires equal treatment of the failure cases, the structural risks, and the regulatory environment that is reshaping the cost of non-compliance.

Regulatory Exposure — EU AI Act

The EU AI Act is now in full effect. For organisations with EU market exposure — which includes any organisation processing data about EU residents, regardless of headquarters location — the Act creates direct financial risk for non-compliant deployment of AI in designated high-risk categories. Credit decisioning, employment assessment, insurance underwriting, and critical infrastructure management are all high-risk categories under the Act. Penalties reach €35 million or 7% of global annual revenue for the most serious violations — a direct cost of non-compliance that is no longer hypothetical. Organisations that have not completed an AI risk assessment for their existing and planned automation programmes have a regulatory exposure they may not have fully priced.

Agentic AI — The Cancellation Risk

Gartner's projection that more than 40% of agentic AI projects will be cancelled by 2027 should inform project selection. The pattern they identify is consistent with the first-wave RPA failure pattern: organisations select processes for agentic automation that appear complex and valuable, but that contain exception densities and undocumented logic that the autonomous system cannot handle. The agent either makes incorrect decisions that create downstream problems, or escalates exceptions at a rate that consumes more human time than the original process.

The mitigation is the same as it was for RPA: process quality before technology deployment, narrow scope on initial programmes, and a measurement framework that detects performance problems before they compound.

The UI Fragility Problem — Still Present

RPA's historical fragility to interface changes has not been fully resolved by modern intelligent automation platforms. Bots that navigate user interfaces remain vulnerable to application updates, and the maintenance overhead of a large bot estate is a real cost that many programme business cases underestimate. Organisations moving from UI-based to API-based integration where the option exists — and building towards more robust system-to-system integration through platforms like MuleSoft or Azure Integration Services — are building more durable automation programmes than those that continue to rely primarily on UI automation.

Applying This by Function: Specific First Steps

The evidence base suggests specific, near-term entry points for each major enterprise function. These are not strategic recommendations — they are process-level starting points based on where the research shows the highest consistency of positive outcomes.

FUNCTION	PROVEN APPLICATION	THE PROBLEM IT SOLVES	FIRST STEP
Operations & Supply Chain	Predictive maintenance (Siemens model)	Unplanned downtime, reactive maintenance costs	Identify 2–3 critical assets where failure causes greatest production disruption. Confirm whether current sensor data is being captured in a format accessible to ML modelling.
Finance & Risk	Contract review automation (COiN / Harvey / Luminance)	Legal review hours on standard documents; inconsistent clause interpretation	Map which contract types consume the most lawyer time. Request demos from Harvey or Luminance on your most common contract category.
Customer Experience	AI-assisted agent response (Iron Mountain / Einstein model)	High handle times; agent context-switching; abandonment	If on Salesforce Service Cloud, request a briefing on Einstein AI features within

FUNCTION	PROVEN APPLICATION	THE PROBLEM IT SOLVES	FIRST STEP
			your current licence.
HR & Talent	Conversational HR assistant (IBM AskHR model)	Repetitive HR enquiries consuming staff time	Segment your last 12 months of HR help desk tickets by query type. Identify the top 5 repeating question categories as first automation targets.
IT & Infrastructure	AIOps and intelligent ticket routing (BT / ServiceNow model)	Reactive incident management; high IT support volume	Pull your last 12 months of IT service desk data. Identify the top 10 ticket categories by volume and request a ServiceNow briefing on AI deflection capability.
Legal & Compliance	Document review and clause extraction	Manual review inconsistency;	Conduct a half-day scoping session to map

FUNCTION	PROVEN APPLICATION	THE PROBLEM IT SOLVES	FIRST STEP
		compliance exposure	which document review workflows consume the most time and have the most consistent structure.
Finance Operations	Invoice processing and AP automation	Manual keying errors; slow approval cycles; duplicate payment risk	Request your current AP error rate data and cycle time data as a baseline before any vendor briefing. The measurement baseline is the starting point.
Healthcare (Admin)	AI ambient scribing and RCM automation	Documentation time; claim denial rates	Pilot ambient scribing with 2–3 willing clinicians and measure documentation time and denial rate change over a 90-day period.

The table is deliberately operational rather than aspirational. Each entry reflects a well-evidenced application with named company results, not a theoretical use case.

What to Watch in the Next 12–18 Months

Platform-embedded AI agents will arrive whether or not you planned for them. Salesforce Agentforce, ServiceNow Now Assist, Microsoft Copilot Studio, and SAP Joule are on active product roadmaps. At your next platform renewal negotiation, AI capability is a pricing variable — understanding which features are included at which licence tiers is a material commercial question, not a technical one.

The EU AI Act compliance window is closing. Organisations with EU exposure that have not completed an AI risk assessment under the Act's high-risk category framework are accruing compliance debt. This is not a technology project. It is a legal and governance project, and it belongs on a 2026 internal agenda regardless of automation maturity.

Workforce redesign is the next constraint. Deloitte's projection that 36% of organisations expect more than 10% of jobs to be fully automated within one year (Deloitte, 2025) suggests that workforce transition planning — role redesign, retraining pathways, communication frameworks — is moving from a future risk to a present-tense operational requirement. Organisations that treat this as an HR problem rather than a leadership problem will find the programmes that worked technically fail at the organisational adoption stage.

Measurement infrastructure is becoming a competitive asset. McKinsey identifies programme measurement capability as a differentiator between high and average automation performers. Organisations investing now in the data infrastructure and governance cadences required to measure automation

effectiveness at scale will have a faster iteration cycle on future programmes than those still debating whether measurement is worth the effort.

Suggested Internal Next Steps

1. Commission a process diagnostic before the next automation procurement decision.

McKinsey's research is unambiguous: organisations that redesign processes before deploying automation are three times more likely to achieve material EBIT impact. Commission a half-day session with your Operations and Finance leads to map the five highest-volume, most repetitive processes each function manages. Use Celonis process mining, your ERP's native analytics, or a structured facilitated session to document where those processes deviate, loop, or generate exceptions. This is prerequisite work — not a nice-to-have.

2. Audit your existing AI and automation licences before the next renewal cycle.

Request briefings from your Salesforce, ServiceNow, Microsoft 365, and SAP account teams specifically on AI features available within your current licence agreements. Many organisations are paying for capability they have not activated. Iron Mountain and ASML both achieved material results by activating AI features within platforms they already owned. Activation of existing tools is lower-cost and lower-risk than new platform procurement.

3. Assign measurement ownership to a named person.

Gartner found fewer than 20% of companies excel at measuring hyperautomation effectiveness. Assign one named person — in IT, finance, or a dedicated automation function — to build and maintain a measurement register for all active automation programmes. The register should include: the baseline captured before deployment; the current performance against that baseline; a defined threshold for intervention; and a documented decision history. Without

measurement, you cannot make evidence-based decisions about where to continue investing, where to scale, or where to stop.

4. Conduct an EU AI Act risk assessment for programmes with EU exposure.

If your organisation has EU market exposure, use the Act's risk classification framework to assess your current and planned automation programmes. High-risk applications in credit, employment, insurance, and critical infrastructure require specific technical and governance standards. This assessment should be completed in 2026, not deferred to a future compliance cycle.

5. Design a workforce transition framework alongside your next major automation programme.

Before launching any automation programme that will materially change the work of a function, design the workforce transition plan in parallel with the technical design. Document which roles change, how they change, what the retraining pathway is, and how affected staff will be communicated with throughout the programme. Deloitte's research is consistent with case study evidence: programmes that invest in this have higher adoption rates, higher quality outputs, and faster scale.

References

AKASA. (2024). *Methodist Health System case study*. AKASA. https://akasa.com/resources/case-studies/methodist-health-system/

Accelirate. (2025). *How AI agents-powered RPA is optimising revenue cycle management in healthcare*. Accelirate. https://www.accelirate.com/rpa-revenue-cycle-management-healthcare/

Bank of America. (2025, April). *AI adoption by BofA's global workforce improves productivity, client service*. Bank of America Newsroom. https://newsroom.bankofamerica.com/content/newsroom/press-releases/2025/04/ai-adoption-by-bofa-s-global-workforce-improves-productivity–cl.html

Bank of America. (2025, August). *A decade of AI innovation: BofA's virtual assistant Erica surpasses 3 billion client interactions*. Bank of America Newsroom. https://newsroom.bankofamerica.com/content/newsroom/press-releases/2025/08/a-decade-of-ai-innovation–bofa-s-virtual-assistant-erica-surpas.html

Benzinga. (2025). IBM's 'Client Zero' AI strategy to deliver $4.5 billion in annual savings. *Benzinga*. https://www.benzinga.com/markets/earnings/25/07/46594954/ibms-client-zero-ai-strategy-to-deliver-4-5-billion-in-annual-savings

Bloomberg. (2017, February 28). JPMorgan software does in seconds what took lawyers 360,000 hours. *Bloomberg*. https://www.bloomberg.com/news/articles/2017-02-28/jpmorgan-marshals-an-army-of-developers-to-automate-high-finance

BT Group. (2024). *BT doubles down on AIOps with Dynatrace, targets self-healing systems by 2025*. BT Newsroom. https://newsroom.bt.com/bt-doubles-down-on-aiops-with-dynatrace–targets-self-healing-systems-by-2025/

Commure. (2025). *Ambient AI clinical documentation outcomes*. Commure. https://www.commure.com/ambient-ai

Council for Affordable Quality Healthcare. (2024). *2024 CAQH Index: Closing the gap — the industry's opportunity to reduce administrative burden*. CAQH. https://www.caqh.org/insights/caqh-index-report

Deloitte. (2025). *Technology trends 2026: Human-AI collaboration*. Deloitte Insights. https://www.deloitte.com/us/en/insights/topics/technology-management/tech-trends/2026/agentic-ai-strategy.html

Ernst & Young. (2023). *Why robotic process automation projects fail and how to fix them*. EY. https://www.ey.com/en_au/consulting/why-rpa-projects-fail

Forrester Research. (2024, October). *The total economic impact of Microsoft 365 Copilot*. Forrester/Microsoft. https://www.microsoft.com/en-us/microsoft-365/blog/2024/10/16/the-total-economic-impact-of-microsoft-365-copilot/

Forrester Research. (2025). *Predictions 2026: Automation at the crossroads*. Forrester. https://www.forrester.com/blogs/predictions-2026-automation-at-the-crossroads/

Gartner. (2024). *Measuring hyperautomation effectiveness: Why most organisations fail*. Gartner. https://www.gartner.com/en/documents/hyperautomation-measurement

Gartner. (2025, August 26). Gartner predicts 40% of enterprise apps will feature task-specific AI agents by 2026, up from less than 5% in 2025 [Press release].

Gartner. https://www.gartner.com/en/newsroom/press-releases/2025-08-26-gartner-predicts-40-percent-of-enterprise-apps-will-feature-task-specific-ai-agents-by-2026-up-from-less-than-5-percent-in-2025

IBM. (2025). *IBM AskHR case study*. IBM. https://www.ibm.com/case-studies/ibm-askhr

IBM. (2025). *IBM Client Zero case study*. IBM. https://www.ibm.com/case-studies/ibm-client-zero

Industrial Equipment News. (2024). Smart factories powered by agentic AI: A new era in manufacturing. *IEN*. https://www.ien.com/artificial-intelligence/blog/22957563/smart-factories-powered-by-agentic-ai-a-new-era-in-manufacturing

KLAS Research. (2025). *Where healthcare is investing and betting on AI in 2025*. KLAS. https://engage.klasresearch.com/blog/where-healthcare-is-investing-and-betting-on-ai-in-2025/8955/

klover.ai. (2025). JPMorgan's AI strategy: Chasing AI dominance. *Klover.ai*. https://www.klover.ai/jpmorgan-ai-strategy-chasing-ai-dominance/

McKinsey & Company. (2025). *Aviva: Rewiring the insurance claims journey with AI*. McKinsey & Company. https://www.mckinsey.com/capabilities/tech-and-ai/how-we-help-clients/rewired-in-action/aviva-rewiring-the-insurance-claims-journey-with-ai

McKinsey & Company. (2025, November). *The state of AI 2025: Agents, innovation, and transformation*. McKinsey & Company. https://www.mckinsey.com/capabilities/quantumblack/our-insights/the-state-of-ai

Microsoft. (2025). *Microsoft 365 Copilot customer stories*. Microsoft. https://www.microsoft.com/en-us/microsoft-365/business/copilot-for-microsoft-365

Nike. (2024). *Nike FY2024 annual report: Supply chain and technology investment outcomes*. Nike Investor Relations. https://investors.nike.com/investors/financial-information/annual-reports/

Notable Health. (2025). *AI and automation in revenue cycle management: Must-know trends for 2025*. Notable Health. https://www.notablehealth.com/blog/ai-and-automation-in-revenue-cycle-management-must-know-trends-for-2025

Salesforce. (2024). *Iron Mountain: AI-faster automation* [Customer story]. Salesforce. https://www.salesforce.com/customer-stories/iron-mountain-ai-faster-automation/

Siemens Digital Industries. (2025). *AI in procurement: Outcomes from Siemens' intelligent supplier management*. Siemens. https://www.siemens.com/global/en/company/stories/research-technologies/digitalization/ai-procurement.html

UPS. (2023). *ORION: How UPS uses AI to optimise routes*. UPS. https://about.ups.com/us/en/our-stories/innovation-driven/orion-how-ups-uses-ai-to-optimize-routes.html

---

Internal use only. Prepared by AI Business Expert. 20 March 2026.